

## "Bacon" SQL features

These features are in the `.sql` files on `.../bacon-numbers-slides-and-companion-code/bacon-numbers/` wherever you unzipped `bacon-numbers-slides-and-companion-code.zip`.

### 1. Classic three-table implementation of a many-to-many actors-plays ERD relationship — `01-cr-tables.sql`

Uses `cast_members` intersection table with, for example:

```
cast_members_fk1 foreign key(actor)
references actors(actor)
match full
on delete cascade
on update restrict
```

### 2. Three-table implementation of a many-to-many actors-actors ERD relationship — `02-cr-edges-table-and-proc.sql`

Uses `edges` intersection table with with an `array` attribute: `movies text[]`.

### 3. Language `p/pgsql` procedure to populate the edges table — `02-cr-edges-table-and-proc.sql`

Uses inner self-join on `edges` and `select... array_agg(movie order by movie)` to populate `movies[]`. Encapsulates `delete` and `insert` as atomic business txn.

### 4. Language `p/pgsql` procedure using dynamic SQL DDLs — `05-cr-raw-paths-table.sql`

Creates table with run-time-specified name with standard column structure. Ancillary objects (like a `sequence`) have systematically related created names.

### 5. Language `p/pgsql` table function — `06-cr-list-paths.sql`

Creates an easy-to-read list of paths where one path is displayed as a sequence of traversed nodes. Uses `translate()` built-in function to render the `::text` typecast of `movies` array more readable *declaratively*.

### 6. Use of a recursive CTE to find all the distinct paths in an undirected cyclic graph (encapsulated in language `p/pgsql` proc) — `07-find-paths-naive.sql`

Shows how to build the paths starting with the `array constructor` in the non-recursive term and `array concatenation` in the recursive term. Uses `p/pgsql` procedure `terminal(some_array)` to encapsulate `some_array[cardinality(some_array)]` to get the last element in an array. Uses the array to prevent cycles by comparing the to-be-added element with every element on the path to date using `any(some_array)`. The procedure encapsulates `drop index`, `delete` from to-be-populated table, and `insert` into it.

### 7. Use of `after insert` trigger on paths to populate a column that traces the execution — `08-cr-raw-paths-with-tracing.sql`

A common technique for use at development time to add execution tracing without cluttering the main code.

### 8. Use of a language `p/pgsql` procedure to implement recursive CTE algorithm as SQL-PL/pgSQL hybrid — `09-find-paths-no-pruning.sql`

This is a well-known technique that database developers come to rely on when the use-case can't be met with a single SQL statement.

### 9. Use of the `@>` ("contains") array operator — `10-cr-restrict-to-unq_containing-paths.sql`

This is used to remove shorter paths that are completely contained by longer paths. It's hard to make up a convincing stand-alone demo that shows the value of this operator. Here, it's the other way round. There's a clear problem to be solved. And "contains" meets the need.

### 10. Use of the `\copy` metacommand to ingest data from an o/s file — `13-insert-imdb-data.sql`

Used to ingest real IMBd — a curated subset from a University site that represents a fully connected graph with so many connected actors and movies (maximum Bacon number is six) that the pure SQL approach never finishes. Shows that the SQL-PL/pgSQL hybrid is *essential*.

## "Employee hierarchy" SQL features

This adds these features that are not used in the "Bacon" use case. They are in the *.sql* files on .../bacon-numbers-slides-and-companion-code/employee-hierarchy/ wherever you unzipped *bacon-numbers-slides-and-companion-code.zip*.

### 11. Creates *domain name\_t as text* — *1-cr-table.sql*

This is used for the *emps.name* and *emps.mgr\_name* columns. Constrains the max length to 30 characters. Much easier to change than if the column were simply *varchar(30)*. Also enforces the rule that names are all lower case.

### 12. Unique expression-based, partial index — *1-cr-table.sql*

Enforces the business rule that there is no more than one ultimate manager.

### 13. Self-referential FK constraint (a.k.a. pig's ear) — *1-cr-table.sql*

Together with the unique expression-based, partial index, this enforces the rule that every employee except for the ultimate manager must have exactly one manager—i.e. that the data is a strict rooted tree (a.k.a. hierarchy).

### 14. Use of *order by... nulls first* — *1-cr-table.sql*

Lists employees in manager order with the ultimate manager first.